Relational data

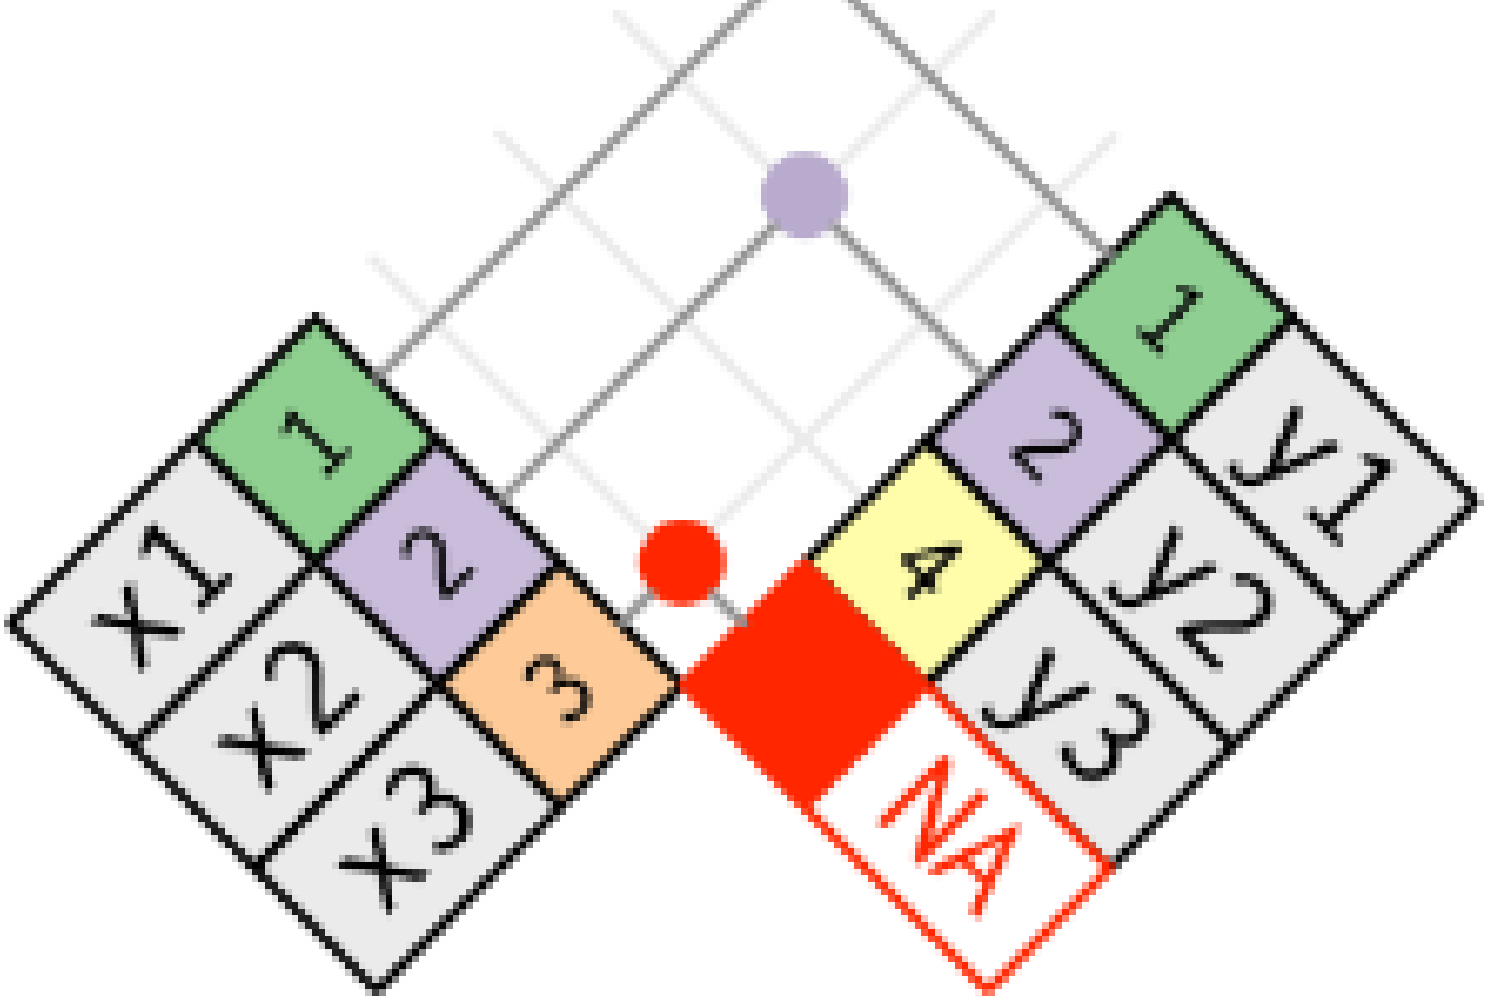

Image: R for Data Science

Relational data

It's rare to find all the data you need for an analysis in a single table.

Typically, you'll have to link two (or more) tables together by matching on common "key" variable(s).

We use joins in SQL *or* R *(or VLOOKUP in* Excel*)*

Relational data

Here, we'll focus on left (outer) joins.

The syntax is similar for other types of join.

left_join

```
table_1 %>%
    left_join(table_2, by = "x")
```

Keep structure of table_1

...and match to observations in table_2

"key" variable (common to both tables)

Relational Data

We're going to join two tables - one with cases of tuberculosis by country, one with population by country.

From this new table we can derive a rate.

cases

country	year	cases
A	1999	5
B	1999	9

pop

country	year	pop
B	1999	500
A	1999	3000

Please Import

tb_cases.csv

and

tb_pop.csv

left_join

Keep the original structure of the tb_cases data frame

```
tb_cases %>%
  left_join(tb_pop, by = "country")
```

...then match to rows in tb_pop

based on "country" value

Duplicates!

```
tb_cases %>%
  left_join(tb_pop, by = "country")
```

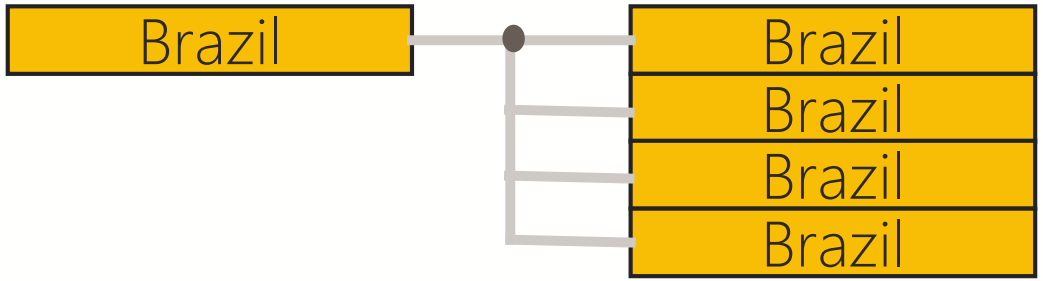

For every value of Brazil in tb_cases, there are 4 in tb_pop...

Join on multiple rows

```
tb_cases %>%
  left_join(tb_pop, by = c("country" , "year"))
```

match on two variables

c stands for 'combine'

Joining with different names

If two tables have different names for same variable:

```
tb_cases %>%
   left_join(bad_names,
   by = c("country" = "Place" , "year" = "Yr"))
```

imaginary table

name in cases

name in bad_names

Some other dplyr joins

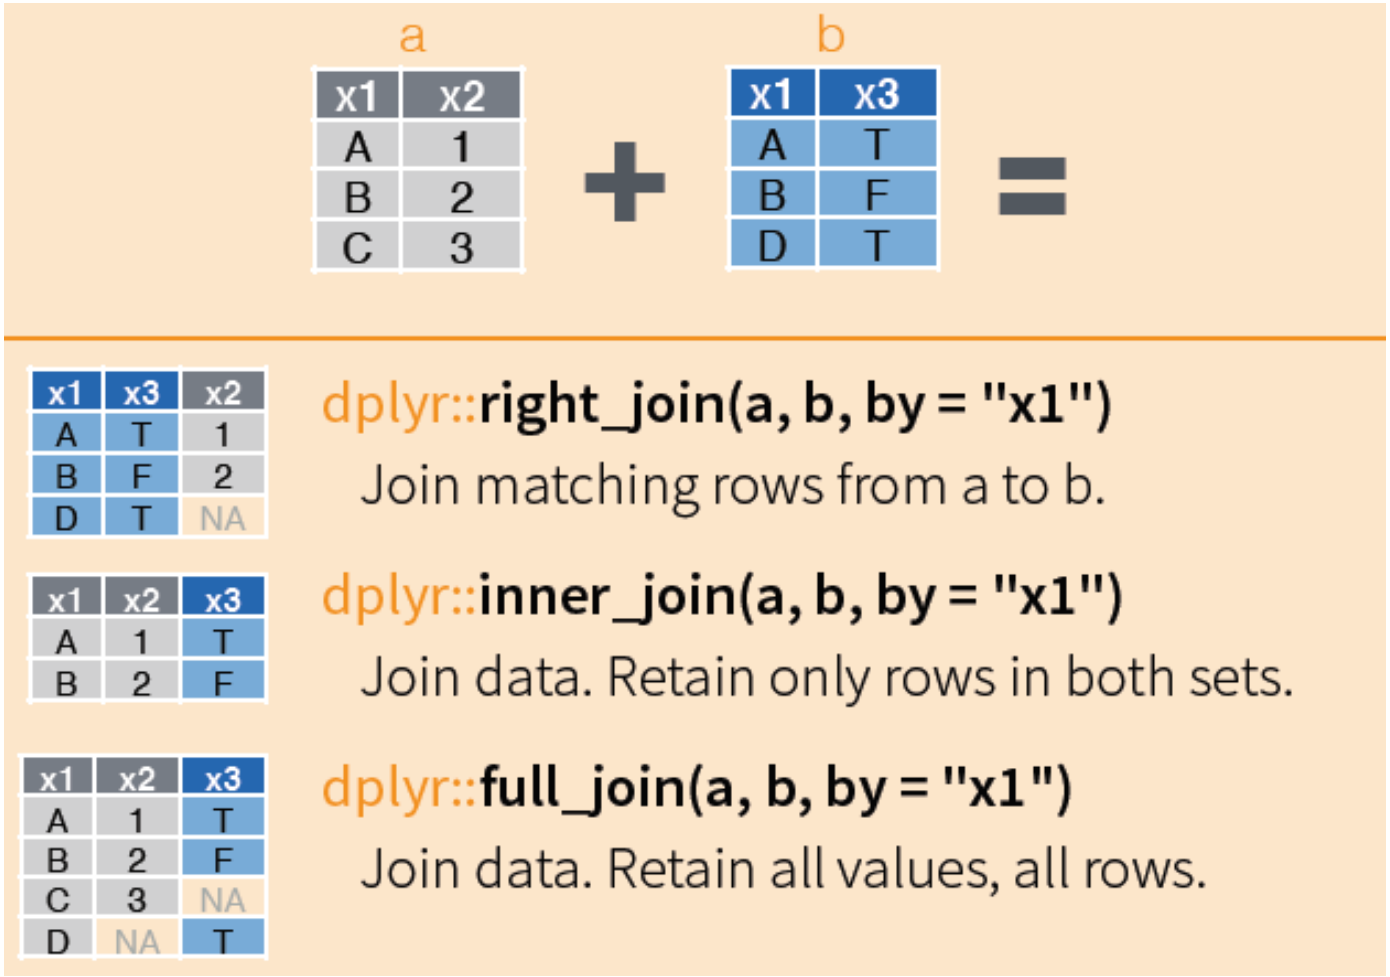

Image taken from: https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf

This work is licensed as

End